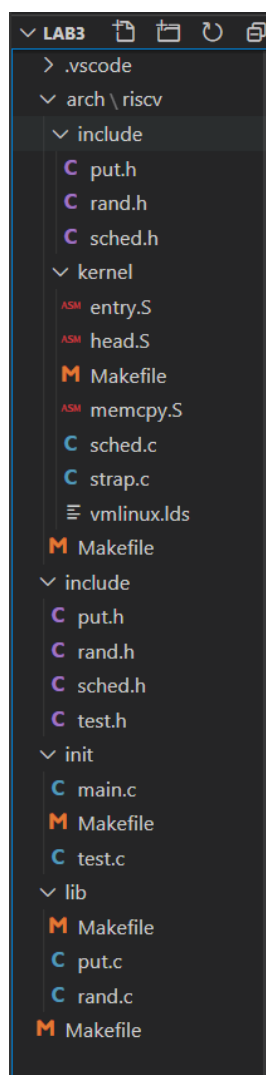


4.1

建立映射

```
cd 'C:\Users\administrater\Desktop\Archived Courses\OS\docker_vol\lab3'  
docker run -it -v ${pwd}:/home/oslab/lab3 -u oslab -w /home/oslab 6014 /bin/bash
```

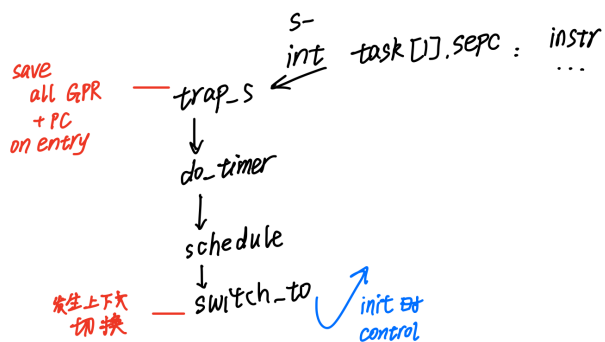
目录结构



存在几个.h文件拷贝多份的现象，不是很好。我认为部分原因是代码层次划分的不是很清楚，比如sched.h里的函数在init文件夹中也需要被调用用于初始化。

4.4 进程调度 调用分析

由于进程调度在实现上，是操作系统在`Supervisor Interrupt Handler`中“干了私活”，因此我们从lab3打下的基础，`entry.S`中的`trap_s`开始跟踪。最终的设计是这样的：



先考虑一般切换，初始化的之后再考虑

entry.S中的trap_s

```

# tracking from here. Maintain stack balance!
trap_s:
#! sp unchanged
# push GPR
pusha
# push sepc
csrr a1, sepc
addi
sd a1

jal do_timer
ecall
# pop sepc
ld a1
addi
csrw sepc, a1
# pop GPR
popa
sret

```

截图中 细节内容/调试输出 模糊处理，下同。

- 等到do_timer的层级调用返回，如果发生了进程切换，sp就变成下一个task的了，pop出来的也是一个task的状态
- 没有内核栈，trap_s及其后调用的函数复用被trap进程的栈。不是很好，但简单。
- 接下来要注意保持栈平衡。尤其是pop时

sched.c中的do_timer

```

#ifdef SJF
void do_timer(void)
{
    current->counter--;

    // PREEMPT_DISABLE
    if (current->counter == 0)
    {
        schedule();
    }
}
#endif

#ifdef PRIORITY
void do_timer(void)
{
    current->counter--;
    if (current->counter == 0)
    {
        current->counter = COUNTER_INIT[current->pid];
    }
    // PREEMPT_ENABLE
    schedule();
}
#endif

```

照实验指导 4.4.3翻译

sched.c中的schedule

照实验指导 4.4.4翻译，分离了重复的调试输出，应该适合阅读。

sched.c中的switch_to

- 这里是当时实现的时候最头疼的地方。但最后做出来，挺naive的，路子挺野

```
// save current
asm(
    "sd ra, 0(%0);\n"
    "sd sp, 8(%0);\n"
    "sd s0,16(%0);\n"
    "sd s1,24(%0);\n"
    "sd s2,32(%0);\n"
    "sd s3,40(%0);\n"
    "sd s4,48(%0);\n"
    "sd s5,56(%0);\n"
    "sd s6,64(%0);\n"
    "sd s7,72(%0);\n"
    "sd s8,80(%0);\n"
    "sd s9,88(%0);\n"
    "sd s10,96(%0);\n"
    "sd s11,104(%0);\n"
    :
    : "r"(&current->thread)
    : "memory");
// load next
asm(
    "ld ra, 0(%0);\n"
    "ld sp, 8(%0);\n"
    "ld s0,16(%0);\n"
    "ld s1,24(%0);\n"
    "ld s2,32(%0);\n"
    "ld s3,40(%0);\n"
    "ld s4,48(%0);\n"
    "ld s5,56(%0);\n"
    "ld s6,64(%0);\n"
    "ld s7,72(%0);\n"
    "ld s8,80(%0);\n"
    "ld s9,88(%0);\n"
    "ld s10,96(%0);\n"
    "ld s11,104(%0);\n"
    :
    : "r"(&next->thread)
    :);
// gcc shouldn't know this
// 环境恶劣
current = next;
```

从gcc的角度看，我们其实是破坏了C函数的调用规范，把s*寄存器都摧毁了，还动了栈顶指针，非常恶劣。但是我们是在进行系统编程。除了帮我们算两个task->thread的偏移之外，gcc绝不能插手这一切。要小心gcc自作聪明。

dump最后的可执行文件vmlinux，确认符合预期。

```

0000000080000a00 <switch_to>:
80000a00: 00000717          auipc   a4, 0x0
80000a04: 2d070713          addi    a4, a4, 720 # 80000cd0 <current>
80000a08: 00073783          ld      a5, 0(a4)
80000a0c: 08a78063          beq     a5, a0, 80000a8c <switch_to+0x8c>
80000a10: 02878793          addi    a5, a5, 40
80000a14: 0017b023          sd      ra, 0(a5)
80000a18: 0027b423          sd      sp, 8(a5)
80000a1c: 0087b823          sd      s0, 16(a5)
80000a20: 0097bc23          sd      s1, 24(a5)
80000a24: 0327b023          sd      s2, 32(a5)
80000a28: 0337b423          sd      s3, 40(a5)
80000a2c: 0347b823          sd      s4, 48(a5)
80000a30: 0357bc23          sd      s5, 56(a5)
80000a34: 0567b023          sd      s6, 64(a5)
80000a38: 0577b423          sd      s7, 72(a5)
80000a3c: 0587b823          sd      s8, 80(a5)
80000a40: 0597bc23          sd      s9, 88(a5)
80000a44: 07a7b023          sd      s10, 96(a5)
80000a48: 07b7b423          sd      s11, 104(a5)
80000a4c: 02850793          addi    a5, a0, 40
80000a50: 0007b083          ld      ra, 0(a5)
80000a54: 0087b103          ld      sp, 8(a5)
80000a58: 0107b403          ld      s0, 16(a5)
80000a5c: 0187b483          ld      s1, 24(a5)
80000a60: 0207b903          ld      s2, 32(a5)
80000a64: 0287b983          ld      s3, 40(a5)
80000a68: 0307ba03          ld      s4, 48(a5)
80000a6c: 0387ba83          ld      s5, 56(a5)
80000a70: 0407bb03          ld      s6, 64(a5)
80000a74: 0487bb83          ld      s7, 72(a5)
80000a78: 0507bc03          ld      s8, 80(a5)
80000a7c: 0587bc83          ld      s9, 88(a5)
80000a80: 0607bd03          ld      s10, 96(a5)
80000a84: 0687bd83          ld      s11, 104(a5)
80000a88: 00a73023          sd      a0, 0(a4)
80000a8c: 00008067          ret

```

初看lab3指导的时候，我其实有个困惑：

有点诡异，现在是在进行中断编程，被打断进程的所有寄存器都应该保存才对（包括 temporary）

后来自己做做，想明白，这其实是保护C函数的调用规范，保callee-saved registers，剩下的工作，由切换到下一个task的栈之后，函数逐级返回，退栈时完成。尤其是trap_s的最后部分，会恢复所有的通用寄存器和epc。

sched.c中的task_init

特别注意 栈不平 的问题。第一次切换到子进程时，由于子进程根本就没有被reschedule过，它的sp指在栈底，从switch_to逐级走正常的C函数返回流程，栈就会underflow.

我的解决方法：把子进程的thread.ra偷偷指到这个函数：

```

void task_epc_init()
{
    // led control here with RET;
    asm("csrw sepc, %0;\n\
        ecall;\n\
        sret;"
        :
        : "r"(dead_loop)
        :);
}

```

初始化epc为task将要执行的函数，然后直接接trap_s的sret返回，跳过中间层级。

4.5 测试

Priority 示例

```

ZJU OS LAB 3          3180103008
[PID = 1] Process Created Successfully! counter = 7 priority = 5
[PID = 2] Process Created Successfully! counter = 6 priority = 5
[PID = 3] Process Created Successfully! counter = 5 priority = 5
[PID = 4] Process Created Successfully! counter = 4 priority = 5
[!] Switch from task 0 to task 4, prio: 5, counter: 4
tasks' priority changed
[PID = 1] counter = 7 priority = 1
[PID = 2] counter = 6 priority = 4
[PID = 3] counter = 5 priority = 5
[PID = 4] counter = 4 priority = 4
[!] Switch from task 4 to task 1, prio: 1, counter: 7
tasks' priority changed
[PID = 1] counter = 7 priority = 5
[PID = 2] counter = 6 priority = 5
[PID = 3] counter = 5 priority = 5
[PID = 4] counter = 3 priority = 2
[!] Switch from task 1 to task 4, prio: 2, counter: 3
tasks' priority changed
[PID = 1] counter = 6 priority = 4
[PID = 2] counter = 6 priority = 4
[PID = 3] counter = 5 priority = 4
[PID = 4] counter = 3 priority = 5
[!] Switch from task 4 to task 3, prio: 4, counter: 5
tasks' priority changed
[PID = 1] counter = 6 priority = 5
[PID = 2] counter = 6 priority = 5
[PID = 3] counter = 5 priority = 4
[PID = 4] counter = 2 priority = 2
[!] Switch from task 3 to task 4, prio: 2, counter: 2
tasks' priority changed
[PID = 1] counter = 6 priority = 5
[PID = 2] counter = 6 priority = 3
[PID = 3] counter = 4 priority = 3
[PID = 4] counter = 2 priority = 4
[!] Switch from task 4 to task 3, prio: 3, counter: 4
tasks' priority changed
[PID = 1] counter = 6 priority = 1
[PID = 2] counter = 6 priority = 3
[PID = 3] counter = 4 priority = 5
[PID = 4] counter = 1 priority = 3
[!] Switch from task 3 to task 1, prio: 1, counter: 6
tasks' priority changed

```

体会 & references

- LAST_TASK定义的不好
- 实验指导 task[0] counter被初始化为0，挺诡异的。变成了 先有鸡还是先有蛋 的问题
- 用于测试的task本身代码过于简单，无法从task本身的运行输出，区分获得控制(pc)的是哪个task。
建议改进

汇编与C的互相访问比较困难。

- C访问寄存器 gcc内联汇编 <https://blog.csdn.net/lwx62/article/details/82796364>
- 汇编访问C结构体 定位困难
- 相比我的switch_to，Linux源代码采用宏和shell脚本预处理的方法结合，避免hard code，提高可维护性。 <https://blog.csdn.net/p0x1307/article/details/44492457>
- 做子进程fork的时候，gcc由我的代码推断出一个memcpy，并因为没有实现而报了个错。后来，参考 <https://elixir.bootlin.com/linux/v4.20/source> 提供memcpy实现